

# Basic Operations

This lecture will cover some basic operations with Spark DataFrames.

We will play around with some stock data from Apple.

```
In [1]: from pyspark.sql import SparkSession

In [2]: # May take awhile Locally
spark = SparkSession.builder.appName("Operations").getOrCreate()

In [26]: # Let Spark know about the header and infer the Schema types!
df = spark.read.csv('appl_stock.csv',inferSchema=True,header=True)

In [28]: df.printSchema()

root
 |-- Date: timestamp (nullable = true)
 |-- Open: double (nullable = true)
 |-- High: double (nullable = true)
 |-- Low: double (nullable = true)
 |-- Close: double (nullable = true)
 |-- Volume: integer (nullable = true)
 |-- Adj Close: double (nullable = true)
```

# Filtering Data

A large part of working with DataFrames is the ability to quickly filter out data based on conditions. Spark DataFrames are built on top of the Spark SQL platform, which means that is you already know SQL, you can quickly and easily grab that data using SQL commands, or using the DataFrame methods (which is what we focus on in this course).

```
In [31]: # Using SQL
df.filter("Close<500").show()
```

Date	Open	High	Low	Close	Volume
2010-01-04 00:00:...	213.429998	214.499996	212.380000	214.009998	12343240
2010-01-05 00:00:...	214.599998	215.589994	213.249994	214.379993	15047620
2010-01-06 00:00:...	214.379993	215.23	210.750004	210.969995	13804000
2010-01-07 00:00:...	211.75	212.000006	209.050005	210.58	11928280
2010-01-08 00:00:...	210.299994	212.000006	209.060005	211.980004	11190270
2010-01-11 00:00:...	212.799997	213.000002	208.450005	210.110002	11555740
2010-01-12 00:00:...	209.189994	209.769995	206.419998	207.720001	14861490
2010-01-13 00:00:...	207.870005	210.929995	204.099998	210.650002	15147300
2010-01-14 00:00:...	210.110002	210.459997	209.020004	209.43	10822350
2010-01-15 00:00:...	210.929995	211.599997	205.869999	205.93	14851690
2010-01-19 00:00:...	208.330002	215.189999	207.240004	215.039995	18250190
2010-01-20 00:00:...	214.910006	215.549994	209.500002	211.73	15303820
2010-01-21 00:00:...	212.079994	213.309995	207.210003	208.069996	15203860
2010-01-22 00:00:...	206.780006	207.499996	197.16	197.75	22044190
2010-01-25 00:00:...	202.510002	204.699999	200.190002	203.070002	26642490
2010-01-26 00:00:...	205.950001	213.710005	202.580004	205.940001	46677750

```
|2010-01-27 00:00:...|      206.849995|      210.58|      199.530001|      207.880005|43064210
0|26.932840000000002|
|2010-01-28 00:00:...|      204.930004|      205.500004|      198.699995|      199.289995|29337560
0|25.819922000000002|
|2010-01-29 00:00:...|      201.079996|      202.199995|      190.250002|      192.060003|31148810
0|      24.883208|
|2010-02-01 00:00:...|192.36999699999998|      196.0|191.29999899999999|      194.729998|18746910
0|      25.229131|
+-----+-----+-----+-----+-----+
+-----+
only showing top 20 rows
```

```
In [35]: # Using SQL with .select()
df.filter("Close<500").select('Open').show()
```

```
+-----+
|      Open|
+-----+
|      213.429998|
|      214.599998|
|      214.379993|
|      211.75|
|      210.299994|
|212.799997000000002|
|209.18999499999998|
|      207.870005|
|210.11000299999998|
|210.929995000000002|
|      208.330002|
|      214.910006|
|      212.079994|
|206.78000600000001|
|202.51000200000001|
|205.95000100000001|
|      206.849995|
|      204.930004|
|      201.079996|
|192.36999699999998|
+-----+
only showing top 20 rows
```

```
In [36]: # Using SQL with .select()
df.filter("Close<500").select(['Open', 'Close']).show()
```

```
+-----+-----+
|      Open|      Close|
+-----+-----+
|      213.429998|      214.009998|
|      214.599998|      214.379993|
|      214.379993|      210.969995|
|      211.75|      210.58|
|      210.299994|211.98000499999998|
|212.799997000000002|210.11000299999998|
|209.18999499999998|      207.720001|
|      207.870005|      210.650002|
|210.11000299999998|      209.43|
|210.929995000000002|      205.93|
|      208.330002|      215.039995|
|      214.910006|      211.73|
|      212.079994|      208.069996|
|206.78000600000001|      197.75|
|202.51000200000001|      203.070002|
|205.95000100000001|      205.940001|
|      206.849995|      207.880005|
|      204.930004|      199.289995|
|      201.079996|      192.060003|
|192.36999699999998|      194.729998|
+-----+-----+
only showing top 20 rows
```

Using normal python comparison operators is another way to do this, they will look very similar to SQL operators, except you need to make sure you are calling the entire column within the dataframe, using the format: `df["column name"]`

Let's see some examples:

```
In [38]: df.filter(df["Close"] < 200).show()
```

```
In [39]: df.filter(df['Close'] < 200).show()
```

```
+-----+-----+-----+-----+-----+-----+
+-----+
|          Date|          Open|          High|          Low|          Close|  Volum
e|          Adj Close|
+-----+-----+-----+-----+-----+-----+
+-----+
|2010-01-22 00:00:...|206.78000600000001|          207.499996|          197.16|          197.75|22044190
0|          25.620401|
|2010-01-28 00:00:...|          204.930004|          205.500004|          198.699995|          199.289995|29337560
0|25.819922000000002|
|2010-01-29 00:00:...|          201.079996|          202.199995|          190.250002|          192.060003|31148810
0|          24.883208|
|2010-02-01 00:00:...|192.36999699999998|          196.0|191.29999899999999|          194.729998|18746910
0|          25.229131|
|2010-02-02 00:00:...|          195.909998|          196.319994|193.37999299999998|          195.859997|17458560
0|25.375532999999997|
|2010-02-03 00:00:...|          195.169994|          200.200003|          194.420004|          199.229994|15383200
0|25.812148999999998|
|2010-02-04 00:00:...|          196.730003|          198.370001|          191.570005|          192.050003|18941300
0|          24.881912|
|2010-02-05 00:00:...|192.63000300000002|          196.0|          190.850002|          195.460001|21257670
0|25.323710000000002|
|2010-02-08 00:00:...|          195.690006|197.88000300000002|          193.999994|194.11999699999998|11956770
0|          25.1501|
|2010-02-09 00:00:...|          196.419996|          197.499994|          194.749998|196.19000400000002|15822170
0|          25.418289|
|2010-02-10 00:00:...|          195.889997|          196.6|          194.26|195.12000700000002| 9259040
0|          25.27966|
|2010-02-11 00:00:...|          194.880001|          199.750006|194.05999599999998|          198.669994|13758640
0|          25.739595|
|2010-02-23 00:00:...|          199.999998|          201.330002|          195.709993|          197.059998|14377370
0|          25.531005|
|2014-06-09 00:00:...|          92.699997|          93.879997|          91.75|          93.699997| 7541500
0|          88.906324|
|2014-06-10 00:00:...|          94.730003|          95.050003|          93.57|          94.25| 6277700
0|          89.428189|
|2014-06-11 00:00:...|          94.129997|          94.760002|          93.470001|          93.860001| 4568100
0|          89.058142|
|2014-06-12 00:00:...|          94.040001|          94.120003|          91.900002|          92.290001| 5474900
0|          87.568463|
|2014-06-13 00:00:...|          92.199997|          92.440002|          90.879997|          91.279999| 5452500
0|          86.610132|
|2014-06-16 00:00:...|          91.510002|          92.75|          91.449997|          92.199997| 3556100
0|          87.483064|
|2014-06-17 00:00:...|          92.309998|          92.699997|          91.800003| 92.08000200000001| 2972600
0| 87.36920699999999|
+-----+-----+-----+-----+-----+-----+
+-----+
only showing top 20 rows
```

```
In [39]: # Will produce an error, make sure to read the error!
df.filter(df["Close"] < 200 and df['Open'] > 200).show()
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-39-df4ea1e41a0f> in <module>()
----> 1 df.filter(df["Close"] < 200 and df['Open'] > 200).show()

/usr/local/Cellar/apache-spark/2.0.1/libexec/python/pyspark/sql/column.py in __nonzero__(self)
    425
    426     def __nonzero__(self):
-> 427         raise ValueError("Cannot convert column into bool: please use '&' for 'and', '|' for 'or',
    "
    428                                     "'~' for 'not' when building DataFrame boolean expressions.")
    429     __bool__ = __nonzero__

ValueError: Cannot convert column into bool: please use '&' for 'and', '|' for 'or', '~' for 'not' when bu
ilding DataFrame boolean expressions.
```

```
In [47]: # Make sure to add in the parenthesis separating the statements!
df.filter( (df["Close"] < 200) & (df['Open'] > 200) ).show()
```

```
+-----+-----+-----+-----+-----+-----+
+-----+
|          Date|          Open|          High|          Low|          Close|  Volume|          Adj Close|
+-----+-----+-----+-----+-----+-----+
|2010-01-22 00:00:...|206.78000600000001|          207.499996|          197.16|          197.75|220441900|          25.620401|
|2010-01-28 00:00:...|          204.930004|          205.500004|          198.699995|          199.289995|293375600|25.819922000000002|
|2010-01-29 00:00:...|          201.079996|          202.199995|          190.250002|          192.060003|311488100|          24.883208|
+-----+-----+-----+-----+-----+-----+
only showing top 3 rows
```

```
|2010-01-25 00:00:...|201.079999|202.199999|190.250000|192.000000|211700100|24.883208|
+-----+-----+-----+-----+-----+-----+-----+
```

```
In [49]: # Make sure to add in the parenthesis separating the statements!
df.filter( (df["Close"] < 200) | (df['Open'] > 200) ).show()
```

```
+-----+-----+-----+-----+-----+-----+-----+
+-----+
|      Date|      Open|      High|      Low|      Close|  Volum
e|      Adj Close|
+-----+-----+-----+-----+-----+-----+-----+
+-----+
|2010-01-04 00:00:...|      213.429998|      214.499996|212.38000099999996|      214.009998|12343240
0|      27.727039|
|2010-01-05 00:00:...|      214.599998|      215.589994|      213.249994|      214.379993|15047620
0|27.774976000000002|
|2010-01-06 00:00:...|      214.379993|      215.23|      210.750004|      210.969995|13804000
0|27.333178000000004|
|2010-01-07 00:00:...|      211.75|      212.000006|      209.050005|      210.58|11928280
0|      27.28265|
|2010-01-08 00:00:...|      210.299994|      212.000006|209.06000500000002|211.98000499999998|11190270
0|      27.464034|
|2010-01-11 00:00:...|212.79999700000002|      213.000002|      208.450005|210.11000299999998|11555740
0|      27.221758|
|2010-01-12 00:00:...|209.18999499999998|209.76999500000002|      206.419998|      207.720001|14861490
0|      26.91211|
|2010-01-13 00:00:...|      207.870005|210.92999500000002|      204.099998|      210.650002|15147300
0|      27.29172|
|2010-01-14 00:00:...|210.11000299999998|210.45999700000002|      209.020004|      209.43|10822350
0|      27.133657|
|2010-01-15 00:00:...|210.92999500000002|211.59999700000003|      205.869999|      205.93|14851690
0|26.680197999999997|
|2010-01-19 00:00:...|      208.330002|215.18999900000003|      207.240004|      215.039995|18250190
0|27.860484999999997|
|2010-01-20 00:00:...|      214.910006|      215.549994|      209.500002|      211.73|15303820
0|      27.431644|
|2010-01-21 00:00:...|      212.079994|213.30999599999998|      207.210003|      208.069996|15203860
0|      26.957455|
|2010-01-22 00:00:...|206.78000600000001|      207.499996|      197.16|      197.75|22044190
0|      25.620401|
|2010-01-25 00:00:...|202.51000200000001|      204.699999|      200.190002|      203.070002|26642490
0|26.309658000000002|
|2010-01-26 00:00:...|205.95000100000001|      213.710005|      202.580004|      205.940001|46677750
0|      26.681494|
|2010-01-27 00:00:...|      206.849995|      210.58|      199.530001|      207.880005|43064210
0|26.932840000000002|
|2010-01-28 00:00:...|      204.930004|      205.500004|      198.699995|      199.289995|29337560
0|25.819922000000002|
|2010-01-29 00:00:...|      201.079996|      202.199995|      190.250002|      192.060003|31148810
0|      24.883208|
|2010-02-01 00:00:...|192.36999699999998|      196.0|191.29999899999999|      194.729998|18746910
0|      25.229131|
+-----+-----+-----+-----+-----+-----+-----+
+-----+
```

only showing top 20 rows

```
In [51]: # Make sure to add in the parenthesis separating the statements!
df.filter( (df["Close"] < 200) & ~(df['Open'] < 200) ).show()
```

```
+-----+-----+-----+-----+-----+-----+-----+
+-----+
|      Date|      Open|      High|      Low|      Close|  Volume|      Adj Close|
+-----+-----+-----+-----+-----+-----+-----+
+-----+
|2010-01-22 00:00:...|206.78000600000001|207.499996|      197.16|      197.75|220441900|      25.620401|
|2010-01-28 00:00:...|      204.930004|205.500004|198.699995|199.289995|293375600|25.819922000000002|
|2010-01-29 00:00:...|      201.079996|202.199995|190.250002|192.060003|311488100|      24.883208|
+-----+-----+-----+-----+-----+-----+-----+
+-----+
```

```
In [46]: df.filter(df["Low"] == 197.16).show()
```

```
+-----+-----+-----+-----+-----+-----+-----+
+-----+
|      Date|      Open|      High|      Low|      Close|  Volume|Adj Close|
+-----+-----+-----+-----+-----+-----+-----+
+-----+
|2010-01-22 00:00:...|206.78000600000001|207.499996|197.16|197.75|220441900|25.620401|
+-----+-----+-----+-----+-----+-----+-----+
+-----+
```

```
In [52]: # Collecting results as Python objects
df.filter(df["Low"] == 197.16).collect()
```

```
df.filter(df["Low"] == 197.16).collect()
```

```
Out[52]: [Row(Date=datetime.datetime(2010, 1, 22, 0, 0), Open=206.78000600000001, High=207.499996, Low=197.16, Close=197.75, Volume=220441900, Adj Close=25.620401)]
```

```
In [53]: result = df.filter(df["Low"] == 197.16).collect()
```

```
In [62]: # Note the nested structure returns a nested row object
         type(result[0])
```

```
Out[62]: pyspark.sql.types.Row
```

```
In [65]: row = result[0]
```

Rows can be called to turn into dictionaries

```
In [64]: row.asDict()
```

```
Out[64]: {'Adj Close': 25.620401,
```